

CS 484 - Introduction to Machine Learning

Final Exam

Problem 1:-

Given Data,

Data ID	X_1	X_2	Y	$P(\text{false} X_1, X_2)$	$P(\text{true} X_1, X_2)$
d_1	-4	-2	true	0.08	0.92
d_2	-2	-1	false	0.18	0.82
d_3	0	0	false	0.38	0.62
d_4	3	2	true	0.62	0.38
d_5	1	-1	false	0.82	0.18

Here, we know that

$$P(\text{true} | X_1, X_2) = 1 - P(\text{false} | X_1, X_2)$$

$$\frac{\partial \mathcal{L}}{\partial \omega_i} = (Y - P(\text{true} | X_1, X_2)) \cdot \frac{\partial z}{\partial \omega_i}$$

$$\text{Here, } \omega = \omega_0 + \omega_1 X_1 + \omega_2 X_2$$

$$\frac{\partial z}{\partial \omega_0} = 1 \quad \frac{\partial z}{\partial \omega_1} = X_1 \quad \frac{\partial z}{\partial \omega_2} = X_2$$

d_1 :-

Given,
 $X_1 = -4, X_2 = -2, P(\text{true} | X_1, X_2) = 0.92$

$Y - P(\text{true} | X_1, X_2)$, Here $Y = \text{True}$
 So, $Y = 1$

$$1 - 0.92 = 0.08$$

$$\frac{\partial \text{CLL}}{\partial w_0} = (Y - P(\text{true} | x_1, x_2)) \cdot \frac{\partial z}{\partial w_0}$$

$$= 0.08 \times 1 = 0.08$$

$$\frac{\partial \text{CLL}}{\partial w_1} = (Y - P(\text{true} | x_1, x_2)) \cdot \frac{\partial z}{\partial w_1}$$

$$= 0.08 \times -4 = -0.32$$

$$\frac{\partial \text{CLL}}{\partial w_2} = (Y - P(\text{true} | x_1, x_2)) \cdot \frac{\partial z}{\partial w_2}$$

$$= 0.08 \times -2 = -0.16$$

So, Here we will be applying the above formulas for all data points.

d_z:-

Given,

$$x_1 = -2, x_2 = -1, Y = \text{False}$$

$$\text{So, } Y = 0$$

$$P(\text{True} | x_1, x_2) = 0.82$$

$$Y - P(\text{true} | x_1, x_2) = 0 - 0.82 = -0.82$$

$$\frac{\partial \text{CLL}}{\partial w_0} = -0.82 \times 1 = -0.82$$

$$\frac{\partial \text{CLL}}{\partial w_1} = -0.82 \times (-2) = 1.64$$

$$\frac{\partial \text{CLL}}{\partial w_2} = -0.82 \times (-1) = 0.82$$

d3:-

Given,

$$x_1 = 0, x_2 = 0, Y = \text{False}$$

$$\text{so, } Y = 0$$

$$P(\text{True} | x_1, x_2) = 0.62$$

$$Y - P(\text{True} | x_1, x_2) = 0 - 0.62 = -0.62$$

$$\frac{\partial \text{CLL}}{\partial w_0} = -0.62 \times 1 = -0.62$$

$$\frac{\partial \text{CLL}}{\partial w_1} = -0.62 \times 0 = 0$$

$$\frac{\partial \text{CLL}}{\partial w_2} = -0.62 \times 0 = 0$$

d4:-

Given,

$$x_1 = 3, x_2 = 2, Y = \text{True}, \text{so } Y = 1$$

$$P(\text{True} | x_1, x_2) = 0.38$$

$$Y - P(\text{True} | x_1, x_2) = 1 - 0.38 = 0.62$$

$$\frac{\partial \text{CLL}}{\partial w_0} = 0.62 \times 1 = 0.62$$

$$\frac{\partial \text{LL}}{\partial w_1} = 0.62 \times 3 = 1.86$$

$$\frac{\partial \text{LL}}{\partial w_2} = 0.62 \times 2 = 1.24$$

ds:-

Given,

$$X_1 = 1, X_2 = -1, Y = \text{False}$$

$$\text{so, } Y = 0$$

$$P(\text{True} | X_1, X_2) = 0.18$$

$$Y - P(\text{True} | X_1, X_2) = 0 - 0.18 = -0.18$$

$$\frac{\partial \text{LL}}{\partial w_0} = -0.18 \times 1 = -0.18$$

$$\frac{\partial \text{LL}}{\partial w_1} = -0.18 \times 1 = -0.18$$

$$\frac{\partial \text{LL}}{\partial w_2} = -0.18 \times (-1) = 0.18$$

so, now let us fill the table given with the values we found.

Data Id	X_1	X_2	Y	$P(\text{False} X_1, X_2)$	$P(\text{True} X_1, X_2)$	$\frac{\partial \text{LL}}{\partial w_0}$	$\frac{\partial \text{LL}}{\partial w_1}$	$\frac{\partial \text{LL}}{\partial w_2}$
d1	-4	-2	True	0.08	0.92	0.08	-0.32	-0.16
d2	-2	-1	False	0.18	0.82	-0.82	1.64	0.82
d3	0	0	False	0.38	0.62	-0.62	0	0
d4	3	2	True	0.62	0.38	0.62	1.86	1.24
d5	1	-1	False	0.82	0.18	-0.18	-0.18	0.18

Problem 2 :-

Given data,

$$E = -(1-t) \ln(1-y) - t \ln(y)$$

$$A = -0.7$$

$$B = 0.6$$

$$C = 0.23$$

$$y = 0.8$$

$$f = 0$$

$$\omega_A = 2$$

$$\omega_B = -3$$

$$\omega_C = 4$$

a) Partial gradient of E with respect to ω_C ($\frac{\partial E}{\partial \omega_C}$):

$$\frac{\partial E}{\partial \omega_C} = \frac{\partial E}{\partial y} \cdot \frac{\partial y}{\partial \omega_C}$$

$$\frac{\partial E}{\partial y} = \frac{y-t}{y(1-y)} = \frac{0.8-0}{0.8(1-0.8)} = \frac{0.8}{(0.8)(0.2)} = 5$$

$$\frac{\partial y}{\partial \omega_C} = \sigma'(c) \cdot \frac{\partial c}{\partial \omega_C}$$

$$\sigma'(c) = y(1-y) = 0.8(1-0.8) = 0.8 \times 0.2 = 0.16$$

$$\frac{\partial y}{\partial \omega_C} = 5 \times 0.16 \times 0.23$$

$$\frac{\partial y}{\partial w_c} = 0.184$$

b) Partial gradient of E with respect to w_B ($\frac{\partial E}{\partial w_B}$):

$$\frac{\partial E}{\partial w_B} = \frac{\partial E}{\partial y} \cdot \frac{\partial y}{\partial c} \cdot \frac{\partial c}{\partial w_B}$$

Here, we know that $\frac{\partial E}{\partial y} = 5$

$$\frac{\partial y}{\partial c} = y(1-y) \cdot w_c = 0.64$$

$$\sigma(A) = \tanh(A)$$

$$\sigma'(A) = 1 - \tanh^2(A)$$

$$z'(A) = A \text{ (since } z(A) = w_A * A)$$

$$(y - t) = 0.8 - 0 = 0.8$$

$$\sigma(B) = \text{sigmoid}(B)$$

$$\sigma'(B) = \text{sigmoid}(B) (1 - \text{sigmoid}(B))$$

$$A = -0.7$$

$$z'(B) = B \text{ (since } z(B) = w_B * B)$$

$$(y - t) = 0.8 - 0 = 0.8$$

$$\sigma'(B) = \text{sigmoid}(0.6) (1 - \text{sigmoid}(0.6)) \approx 0.2350$$

$$y'(A) = 1 - \tanh(-0.7)^2 \approx 0.5806$$

$$z'(B) = 0.6$$

$$\frac{\partial E}{\partial w_B} = 5 \times 0.64 \times 0.06216 = 0.1989$$

c) partial gradient of error with respect to w_A ($\partial E / \partial w_A$)

$$\frac{\partial E}{\partial w_A} = (y - t) \times \sigma'(A) \times z'(A)$$

Here, $y = 0.8$

$$t = 0$$

$$A = -0.7$$

$$\sigma(A) = \tanh(A)$$

$$\sigma'(A) = 1 - \tanh(A)^2$$

$$z'(A) = A \text{ (since } z(A) = w_A * A \text{)}$$

$$(y - t) = 0.8 - 0 = 0.8$$

$$\sigma'(A) = 1 - \tanh(-0.7)^2 = 0.5806$$

$$z'(A) = -0.7$$

$$\frac{\partial E}{\partial w_A} = 0.8 (1 - 0.6) \times (-0.7) = -0.168$$

$$\frac{\partial E}{\partial w_A} = 5 \times 0.64 \times (-0.31) \times (-0.168) = 0.167$$

Problem 3:-

Given,

$$P(x|\mu) = \prod_{i=1}^D \mu_i^{x_i} (1-\mu_i)^{(1-x_i)}$$

$$P(x|M, \pi) = \sum_{k=1}^K \pi_k P(x|\mu_k)$$

$$P(x|\mu_k) = \prod_{i=1}^D \mu_{ki}^{x_{ki}} (1-\mu_{ki})^{(1-x_{ki})}$$

1) we know that,

$$E(x) = \int x (P(x|M, \pi)) dx$$

$$E(x) = \int x \cdot \sum_{k=1}^K \pi_k P(x|\mu_k) dx$$

By linearity,

$$E(x) = \sum_{k=1}^K \pi_k \int x \cdot P(x|\mu_k) dx$$

Here,

For each component $P(x|\mu_k)$, $E(x|\mu_k)$ is μ_k

Since each x_i is with parameters μ_{ki}

$$\text{So, } E\{x | \mu_k\} = \mu_k$$

$$E[x] = \sum_{k=1}^K \pi_k \cdot \mu_k$$

Now, $P(x | \mu_k)$ is given and for

$$x = (x_1, x_2, \dots, x_D)$$

$$E\{x | \mu_k\} = \mu_k$$

Hence, it is proved

$$\Rightarrow \text{Now } \Sigma_k = \text{diag}(\mu_{ki}(1 - \mu_{ki}))$$

$$\text{cov}[x] = \sum_{k=1}^K \pi_k (\Sigma_k + \mu_k \mu_k^T) - E[x] E[x]^T$$

$$E[xx^T] = \sum_{k=1}^K \pi_k E[xx^T | \mu_k]$$

For Bernoulli

$$\text{var}(x_i) = \mu_{ki}(1 - \mu_{ki})$$

$$\Sigma_k = \text{diag}(\mu_k(1 - \mu_k))$$

$$E[xx^T | \mu_k] = \Sigma_k + \mu_k \mu_k^T$$

$$E[xx^T] = \sum_{k=1}^K \pi_k (\Sigma_k + \mu_k \mu_k^T)$$

Now, $\text{cov}(x) = E[x x^T] - E[x] E[x]^T$

$$\text{cov}(x) = \sum_{k=1}^K \pi_k (\Sigma_k + \mu_k \mu_k^T) -$$

$$\left(\sum_{i=1}^K \pi_i \mu_i \right) \left(\sum_{k=1}^K \pi_k \mu_k \right)^T$$

It will become,

$$\text{cov}(x) = \sum_{k=1}^K \pi_k (\Sigma_k + \mu_k \mu_k^T) - E[x] E[x]^T$$

Hence, it is proved.

Problem 4:

+ Code + Text

```
import numpy as np
from sklearn.datasets import load_iris, load_breast_cancer, fetch_20newsgroups
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import StandardScaler
from sklearn.exceptions import ConvergenceWarning
import warnings

warnings.filterwarnings("ignore", category=ConvergenceWarning)

def get_dataset(data_name):
    if data_name == "iris":
        dataset = load_iris()
        return dataset.data, dataset.target
    elif data_name == "breast_cancer":
        dataset = load_breast_cancer()
        return dataset.data, dataset.target
    elif data_name == "20newsgroups":
        topics = ['alt.atheism', 'sci.space', 'rec.sport.baseball', 'sci.med']
        data = fetch_20newsgroups(categories=topics)
        vectorizer = TfidfVectorizer(max_features=1000)
        X_transformed = vectorizer.fit_transform(data.data)
        return X_transformed, data.target

def manual_cv(data, labels, estimator, num_folds=10):
    indices = np.arange(len(labels))
    np.random.shuffle(indices)
    partition_size = len(labels) // num_folds
    fold accuracies = []

    for fold_idx in range(num_folds):
        val_start = fold_idx * partition_size
        val_end = (fold_idx + 1) * partition_size if fold_idx < num_folds - 1 else len(labels)
        val_indices = indices[val_start:val_end]
        train_indices = np.concatenate([indices[:val_start], indices[val_end:]])

        train_data, val_data = data[train_indices], data[val_indices]
        train_labels, val_labels = labels[train_indices], labels[val_indices]

        estimator.fit(train_data, train_labels)
        fold accuracies.append(estimator.score(val_data, val_labels))

    return np.mean(fold accuracies)

def analyze_models(data_name):
    data, labels = get_dataset(data_name)
    train_data, test_data, train_labels, test_labels = train_test_split(data, labels, test_size=0.4, random_state=42)

    if isinstance(train_data, np.ndarray):
        scaler = StandardScaler()
        train_data = scaler.fit_transform(train_data)
        test_data = scaler.transform(test_data)

    hyperparameter_C = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
    feature_count = train_data.shape[1]
    layer_options = [int(feature_count * factor) for factor in [0.1, 0.2, 0.5, 1, 2, 5, 10] if int(feature_count * factor) > 0]

    print(f"\nPerforming Cross-Validation on: {data_name.upper()}")

    print("\nLR Performance:")
    top_lr_accuracy = 0
    optimal_lr_param = None
    for reg_param in hyperparameter_C:
        lr_model = LogisticRegression(C=reg_param, max_iter=1000)
        accuracy = manual_cv(train_data, train_labels, lr_model)
        print(f"Hyperparameter Value={reg_param}: Mean Accuracy = {accuracy:.4f}")
        if accuracy > top_lr_accuracy:
            top_lr_accuracy = accuracy
            optimal_lr_param = reg_param

    print("\nSVM with Linear Kernel Performance:")
    best_linear_svm_score = 0
    best_svm_param = None
    for reg_param in hyperparameter_C:
        linear_svm = SVC(kernel="linear", C=reg_param)
        accuracy = manual_cv(train_data, train_labels, linear_svm)
```

```

print(f" hyperparameter_value={reg_param}: Mean Accuracy = {accuracy:.4f}")
if accuracy > best_linear_svm_score:
    best_linear_svm_score = accuracy
    best_svm_param = reg_param

print("\nSVM with RBF Kernel Performance:")
top_rbf_accuracy = 0
best_rbf_param = None
for reg_param in hyperparameter_C:
    rbf_svm = SVC(kernel="rbf", C=reg_param)
    accuracy = manual_cv(train_data, train_labels, rbf_svm)
    print(f" hyperparameter_value={reg_param}: Mean Accuracy = {accuracy:.4f}")
    if accuracy > top_rbf_accuracy:
        top_rbf_accuracy = accuracy
        best_rbf_param = reg_param

print("\nMLP Performance:")
highest_mlp_accuracy = 0
best_layer_config = None
for layer_size in layer_options:
    mlp_model = MLPClassifier(hidden_layer_sizes=(layer_size,), max_iter=1000)
    accuracy = manual_cv(train_data, train_labels, mlp_model)
    print(f" hidden_layer_sizes={layer_size}: Mean Accuracy = {accuracy:.4f}")
    if accuracy > highest_mlp_accuracy:
        highest_mlp_accuracy = accuracy
        best_layer_config = layer_size

print("\nBest Model Summary:")
print(f" L2-LR: Best hyperparameter_value={optimal_lr_param}, Accuracy={top_lr_accuracy:.4f}")
print(f" SVM with Linear Kernel: Best hyperparameter_value={best_svm_param}, Accuracy={best_linear_svm_score:.4f}")
print(f" SVM with RBF Kernel: Best hyperparameter_value={best_rbf_param}, Accuracy={top_rbf_accuracy:.4f}")
print(f" MLP: Best hidden_layer_sizes={best_layer_config}, Accuracy={highest_mlp_accuracy:.4f}")

return {
    "L2-LR": {"Parameter": optimal_lr_param, "Accuracy": top_lr_accuracy},
    "SVM with Linear Kernel": {"Parameter": best_svm_param, "Accuracy": best_linear_svm_score},
    "SVM with RBF Kernel": {"Parameter": best_rbf_param, "Accuracy": top_rbf_accuracy},
    "MLP": {"Layer Size": best_layer_config, "Accuracy": highest_mlp_accuracy},
}

```

```

print(f" SVM with Linear Kernel: Best hyperparameter_value={best_svm_param}, Accuracy={best_linear_svm_score:.4f}")
print(f" SVM with RBF Kernel: Best hyperparameter_value={best_rbf_param}, Accuracy={top_rbf_accuracy:.4f}")
print(f" MLP: Best hidden_layer_sizes={best_layer_config}, Accuracy={highest_mlp_accuracy:.4f}")

return {
    "L2-LR": {"Parameter": optimal_lr_param, "Accuracy": top_lr_accuracy},
    "SVM with Linear Kernel": {"Parameter": best_svm_param, "Accuracy": best_linear_svm_score},
    "SVM with RBF Kernel": {"Parameter": best_rbf_param, "Accuracy": top_rbf_accuracy},
    "MLP": {"Layer Size": best_layer_config, "Accuracy": highest_mlp_accuracy},
}

if __name__ == "__main__":
    datasets_to_analyze = ["iris", "breast_cancer", "20newsgroups"]
    all_results = {}

    for dataset in datasets_to_analyze:
        all_results[dataset] = analyze_models(dataset)

    print("\nOverall Best and Worst Performing Models Across Datasets:")
    for dataset, results in all_results.items():
        best_model = max(results.items(), key=lambda model: model[1]["Accuracy"])
        worst_model = min(results.items(), key=lambda model: model[1]["Accuracy"])
        print(f"{dataset.upper().replace('_', ' ')}:")
        print(f" Best Model = {best_model[0]} with Accuracy = {best_model[1]['Accuracy']:.4f}")
        print(f" Worst Model = {worst_model[0]} with Accuracy = {worst_model[1]['Accuracy']:.4f}")

```



Performing Cross-Validation on: IRIS

L2-LR Performance:

```
hyperparameter_value=0.001: Mean Accuracy = 0.4111
hyperparameter_value=0.01: Mean Accuracy = 0.8111
hyperparameter_value=0.1: Mean Accuracy = 0.8889
hyperparameter_value=1: Mean Accuracy = 0.9222
hyperparameter_value=10: Mean Accuracy = 0.9444
hyperparameter_value=100: Mean Accuracy = 0.9556
hyperparameter_value=1000: Mean Accuracy = 0.9667
```

SVM with Linear Kernel Performance:

```
hyperparameter_value=0.001: Mean Accuracy = 0.3111
hyperparameter_value=0.01: Mean Accuracy = 0.6667
hyperparameter_value=0.1: Mean Accuracy = 0.9444
hyperparameter_value=1: Mean Accuracy = 0.9667
hyperparameter_value=10: Mean Accuracy = 0.9667
hyperparameter_value=100: Mean Accuracy = 0.9667
hyperparameter_value=1000: Mean Accuracy = 0.9556
```

SVM with RBF Kernel Performance:

```
hyperparameter_value=0.001: Mean Accuracy = 0.3333
hyperparameter_value=0.01: Mean Accuracy = 0.2667
hyperparameter_value=0.1: Mean Accuracy = 0.8667
hyperparameter_value=1: Mean Accuracy = 0.9222
hyperparameter_value=10: Mean Accuracy = 0.9333
hyperparameter_value=100: Mean Accuracy = 0.9333
hyperparameter_value=1000: Mean Accuracy = 0.9444
```

MLP Performance:

```
hidden_layer_sizes=2: Mean Accuracy = 0.7333
hidden_layer_sizes=4: Mean Accuracy = 0.9000
hidden_layer_sizes=8: Mean Accuracy = 0.9333
hidden_layer_sizes=20: Mean Accuracy = 0.9222
hidden_layer_sizes=40: Mean Accuracy = 0.9333
```

Best Model Summary:

```
L2-LR: Best hyperparameter_value=1000, Accuracy=0.9667
SVM with Linear Kernel: Best hyperparameter_value=10, Accuracy=0.9667
SVM with RBF Kernel: Best hyperparameter_value=1000, Accuracy=0.9444
MLP: Best hidden_layer_sizes=8, Accuracy=0.9333
```

✓
1h



Performing Cross-Validation on: BREAST_CANCER



L2-LR Performance:

```
hyperparameter_value=0.001: Mean Accuracy = 0.8858
hyperparameter_value=0.01: Mean Accuracy = 0.9413
hyperparameter_value=0.1: Mean Accuracy = 0.9708
hyperparameter_value=1: Mean Accuracy = 0.9736
hyperparameter_value=10: Mean Accuracy = 0.9647
hyperparameter_value=100: Mean Accuracy = 0.9590
hyperparameter_value=1000: Mean Accuracy = 0.9561
```

SVM with Linear Kernel Performance:

```
hyperparameter_value=0.001: Mean Accuracy = 0.9266
hyperparameter_value=0.01: Mean Accuracy = 0.9561
hyperparameter_value=0.1: Mean Accuracy = 0.9647
hyperparameter_value=1: Mean Accuracy = 0.9648
hyperparameter_value=10: Mean Accuracy = 0.9708
hyperparameter_value=100: Mean Accuracy = 0.9471
hyperparameter_value=1000: Mean Accuracy = 0.9354
```

SVM with RBF Kernel Performance:

```
hyperparameter_value=0.001: Mean Accuracy = 0.6128
hyperparameter_value=0.01: Mean Accuracy = 0.6131
hyperparameter_value=0.1: Mean Accuracy = 0.9384
hyperparameter_value=1: Mean Accuracy = 0.9618
hyperparameter_value=10: Mean Accuracy = 0.9679
hyperparameter_value=100: Mean Accuracy = 0.9620
hyperparameter_value=1000: Mean Accuracy = 0.9648
```

MLP Performance:

```
hidden_layer_sizes=3: Mean Accuracy = 0.9588
hidden_layer_sizes=6: Mean Accuracy = 0.9649
hidden_layer_sizes=15: Mean Accuracy = 0.9677
hidden_layer_sizes=30: Mean Accuracy = 0.9706
hidden_layer_sizes=60: Mean Accuracy = 0.9708
hidden_layer_sizes=150: Mean Accuracy = 0.9735
hidden_layer_sizes=300: Mean Accuracy = 0.9706
```

Best Model Summary:

```
L2-LR: Best hyperparameter_value=1, Accuracy=0.9736
SVM with Linear Kernel: Best hyperparameter_value=10, Accuracy=0.9708
SVM with RBF Kernel: Best hyperparameter_value=10, Accuracy=0.9679
MLP: Best hidden_layer_sizes=150, Accuracy=0.9735
```

✓
1h



Performing Cross-Validation on: 20NEWSGROUPS



L2-LR Performance:

```
hyperparameter_value=0.001: Mean Accuracy = 0.3971
hyperparameter_value=0.01: Mean Accuracy = 0.7021
hyperparameter_value=0.1: Mean Accuracy = 0.8882
hyperparameter_value=1: Mean Accuracy = 0.9424
hyperparameter_value=10: Mean Accuracy = 0.9469
hyperparameter_value=100: Mean Accuracy = 0.9536
hyperparameter_value=1000: Mean Accuracy = 0.9507
```

SVM with Linear Kernel Performance:

```
hyperparameter_value=0.001: Mean Accuracy = 0.2291
hyperparameter_value=0.01: Mean Accuracy = 0.2248
hyperparameter_value=0.1: Mean Accuracy = 0.8708
hyperparameter_value=1: Mean Accuracy = 0.9404
hyperparameter_value=10: Mean Accuracy = 0.9449
hyperparameter_value=100: Mean Accuracy = 0.9483
hyperparameter_value=1000: Mean Accuracy = 0.9462
```

SVM with RBF Kernel Performance:

```
hyperparameter_value=0.001: Mean Accuracy = 0.2320
hyperparameter_value=0.01: Mean Accuracy = 0.2830
hyperparameter_value=0.1: Mean Accuracy = 0.7047
hyperparameter_value=1: Mean Accuracy = 0.9396
hyperparameter_value=10: Mean Accuracy = 0.9432
hyperparameter_value=100: Mean Accuracy = 0.9433
hyperparameter_value=1000: Mean Accuracy = 0.9498
```

MLP Performance:

```
hidden_layer_sizes=100: Mean Accuracy = 0.9500
hidden_layer_sizes=200: Mean Accuracy = 0.9549
hidden_layer_sizes=500: Mean Accuracy = 0.9544
hidden_layer_sizes=1000: Mean Accuracy = 0.9572
hidden_layer_sizes=2000: Mean Accuracy = 0.9616
hidden_layer_sizes=5000: Mean Accuracy = 0.9573
hidden_layer_sizes=10000: Mean Accuracy = 0.9521
```

Best Model Summary:

```
L2-LR: Best hyperparameter_value=100, Accuracy=0.9536
SVM with Linear Kernel: Best hyperparameter_value=100, Accuracy=0.9483
SVM with RBF Kernel: Best hyperparameter_value=1000, Accuracy=0.9498
MLP: Best hidden_layer_sizes=2000, Accuracy=0.9616
```

✓
1h



Best Model Summary:



```
L2-LR: Best hyperparameter_value=100, Accuracy=0.9536
SVM with Linear Kernel: Best hyperparameter_value=100, Accuracy=0.9483
SVM with RBF Kernel: Best hyperparameter_value=1000, Accuracy=0.9498
MLP: Best hidden_layer_sizes=2000, Accuracy=0.9616
```

Overall Best and Worst Performing Models Across Datasets:

IRIS:

```
Best Model = SVM with Linear Kernel with Accuracy = 0.9667
Worst Model = MLP with Accuracy = 0.9333
```

BREAST_CANCER:

```
Best Model = L2-LR with Accuracy = 0.9736
Worst Model = SVM with RBF Kernel with Accuracy = 0.9679
```

20NEWSGROUPS:

```
Best Model = MLP with Accuracy = 0.9616
Worst Model = SVM with Linear Kernel with Accuracy = 0.9483
```

Problem 5:

```
✓ [3] import pandas as pd
0s      import numpy as np
      from sklearn.manifold import MDS
      import matplotlib.pyplot as plt

      cities_df = pd.read_csv('/content/cities.csv', sep=';', header=None)
      city_names = cities_df[0].tolist()
      distances = cities_df.iloc[:, 1:25].values.astype(float)

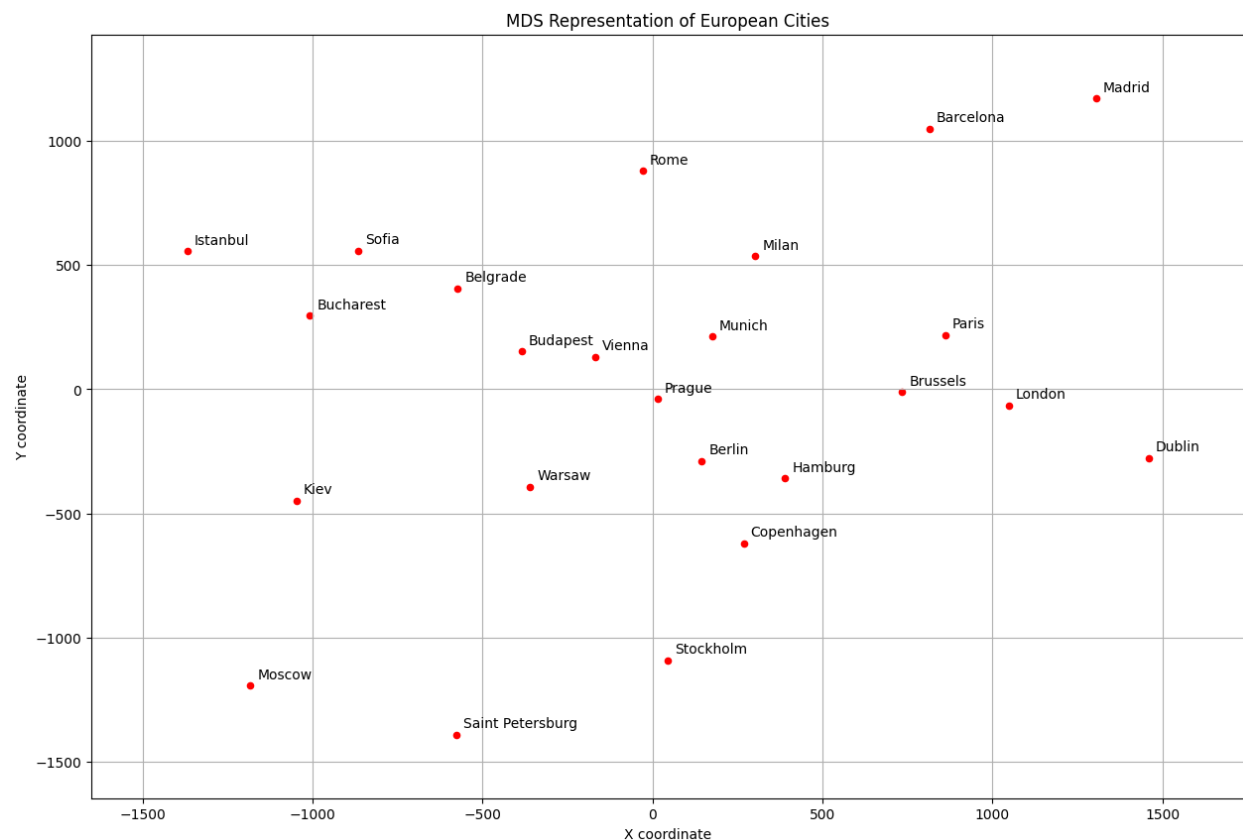
      mds = MDS(n_components=2, dissimilarity='precomputed', random_state=42)
      positions = mds.fit_transform(distances)

      plt.figure(figsize=(15, 10))
      plt.scatter(positions[:, 0], positions[:, 1], color='red', s=20)

      for i, city in enumerate(city_names):
          plt.annotate(city, (positions[i, 0], positions[i, 1]),
                        xytext=(5, 5), textcoords='offset points')

      plt.title('MDS Representation of European Cities')
      plt.xlabel('X coordinate')
      plt.ylabel('Y coordinate')
      plt.grid(True)

      plt.margins(0.1)
      plt.show()
```



Problem 6:

```
✓ [4] import numpy as np
5s import matplotlib.pyplot as plt

def random_walk_vectorized(n_steps, n_simulations=1):
    steps = np.random.choice([-1, 1], size=(n_simulations, n_steps, 2))
    positions = np.cumsum(steps, axis=1)
    return positions

N = 50
plt.figure(figsize=(8, 8))
for i in range(3):
    positions = random_walk_vectorized(N)[0]
    x, y = positions[:, 0], positions[:, 1]
    plt.plot(x, y, label=f"Walk {i+1}")
plt.title(f"Traces of 3 Random Walks with N = {N}")
plt.xlabel("X")
plt.ylabel("Y")
plt.axhline(0, color='black', linewidth=0.5, linestyle='--')
plt.axvline(0, color='black', linewidth=0.5, linestyle='--')
plt.legend()
plt.grid()
plt.show()

def calculate_d_vectorized(n_steps, n_simulations):
    positions = random_walk_vectorized(n_steps, n_simulations)
    final_positions = positions[:, -1, :]
    distances = np.sqrt(np.sum(final_positions**2, axis=1))
    return np.mean(distances)

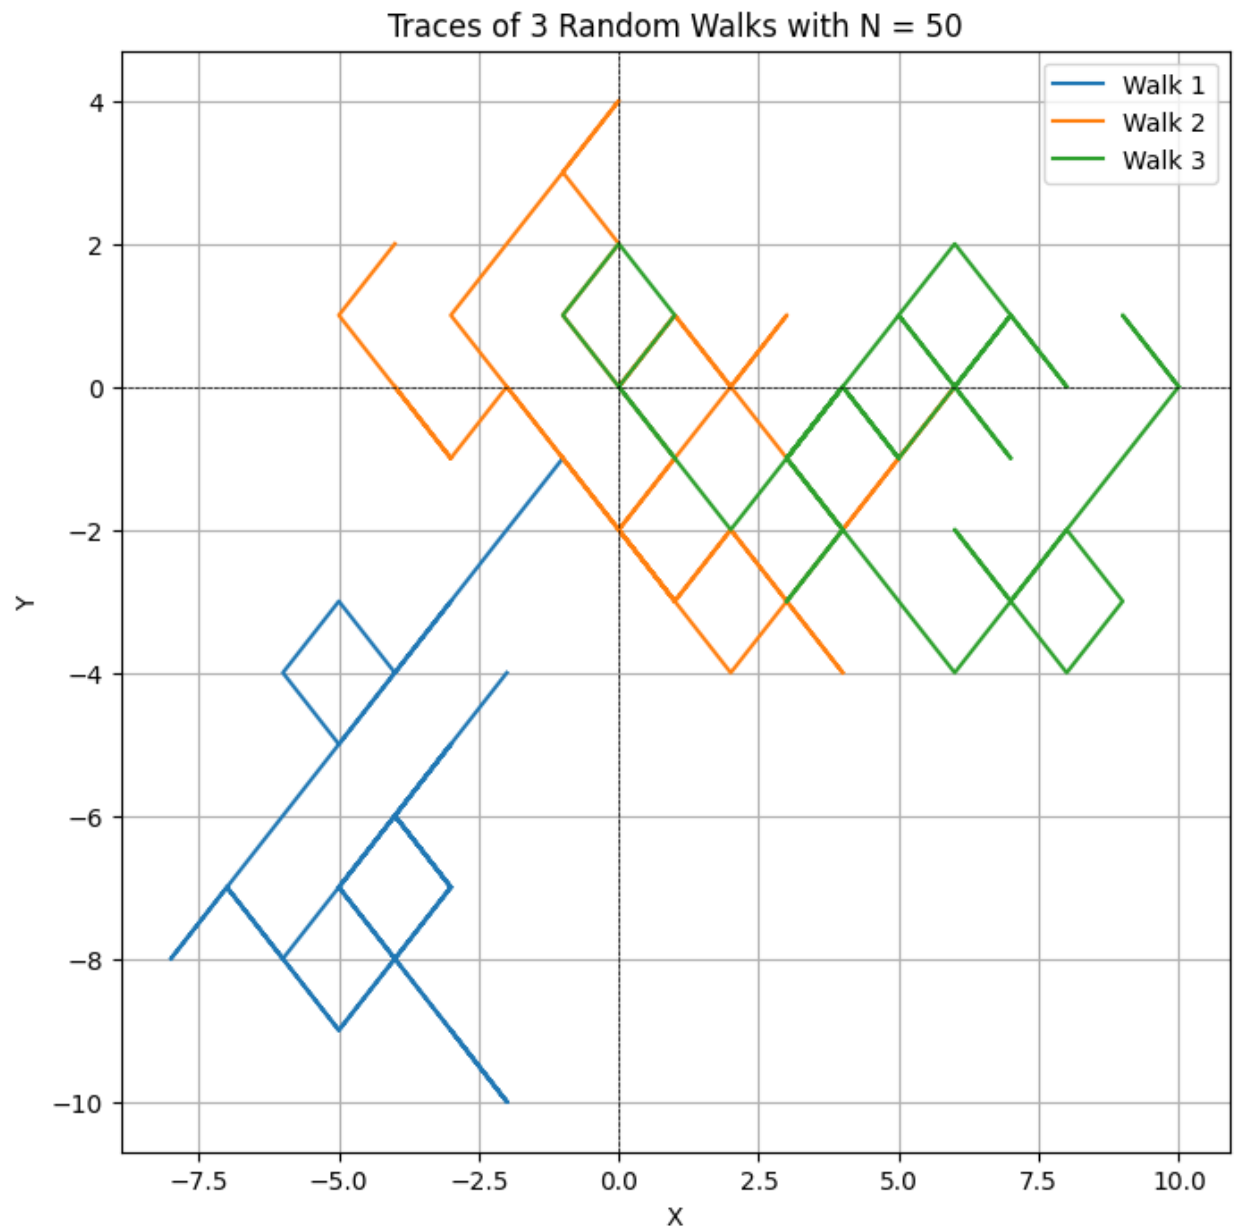
n_simulations = 10000
average_d = calculate_d_vectorized(N, n_simulations)
print(f"Average distance for N = {N} over {n_simulations} simulations: {average_d:.2f}")

n_values = np.arange(10, 501, 10)
average_distances = [calculate_d_vectorized(n, 10000) for n in n_values]

plt.figure(figsize=(8, 6))
plt.plot(np.log(n_values), np.log(average_distances), marker='o', linestyle='--')
plt.title("Scaling of log(d) with log(N)")
```

```
✓ [4] plt.figure(figsize=(8, 6))
5s plt.plot(np.log(n_values), np.log(average_distances), marker='o', linestyle='--')
plt.title("Scaling of log(d) with log(N)")
plt.xlabel("log(N)")
plt.ylabel("log(d)")
plt.grid()

coefficients = np.polyfit(np.log(n_values), np.log(average_distances), 1)
slope = coefficients[0]
print(f"Slope of log-log plot: {slope:.2f}")
plt.show()
```



Average distance for $N = 50$ over 10000 simulations: 8.89
Slope of log-log plot: 0.50

