



Regression vs. Classification





Regression



Classification

Classification Terminology

- **Goal:** Given data points $\{x\} \in R^D$, assign each data to one class C_k ($k= 1, \dots, K$)
- **Decision boundaries:** Input space is divided into regions, whose boundaries are called decision boundaries and **each region corresponds to a class of data**
- **Linearly separable:** Datapoints whose classes can be separated by **linear decision boundaries**
 - mean that decision boundaries are **linear functions of the input x**
 - hence are defined by $(D - 1)$ -dimensional **hyperplanes** within the D -dimensional input space



Classification: Three Different Methods

- **Discriminant models**
 - Given training data, assign each data x to one class C_k via a discriminant function
 - *Do not consider distribution* of the training data
- **Probabilistic discriminant models**
 - Given training data, model the **posterior class distribution** $p(C_k|x)$
 - Use the distribution $p(C_k|x)$ to perform classification for testing data
- **Probabilistic generative models**
 - Given training data, model the **joint (data, class) distribution** $p(x, C_k)$
 - Find class-conditional distribution $p(x|C_k)$ and class prior distribution $p(C_k)$
 - Then use Bayes rule to compute $p(C_k|x) \sim p(x|C_k) p(C_k)$

Discriminant Models

Binary Classification

- The simplest representation of a **linear discriminant function**

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

\mathbf{w} is called a **weight parameter vector**, and w_0 is a **bias**

- An input data \mathbf{x} is classified to


- Class C_1 if $y(x) > 0$
- Class C_2 if $y(x) < 0$

- The decision boundary is defined by

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 = 0$$

Geometry of Linear Discriminant Function

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$



$$y(\mathbf{x}_A) = \mathbf{w}^T \mathbf{x}_A + w_0 = 0$$

$$y(\mathbf{x}_B) = \mathbf{w}^T \mathbf{x}_B + w_0 = 0$$

$$\Rightarrow \mathbf{w}^T (\mathbf{x}_A - \mathbf{x}_B) = 0$$

$$\mathbf{x} = \mathbf{x}_* + r \frac{\mathbf{w}}{\|\mathbf{w}\|_2}$$

$$\mathbf{w}^T \mathbf{x} + w_0 = \mathbf{w}^T \mathbf{x}_* + w_0 + r \frac{\mathbf{w}^T \mathbf{w}}{\|\mathbf{w}\|_2}$$

$$= 0 + r \|\mathbf{w}\|_2$$

$$\Rightarrow r = \frac{\mathbf{w}^T \mathbf{x} + w_0}{\|\mathbf{w}\|_2}$$


$$= \frac{y(\mathbf{x})}{\|\mathbf{w}\|_2}$$

Multi-Class Classification


- Build a **K-class discriminant function** by **combining** a number of **two-class** discriminant functions
 - **One-versus-the-rest classifier**
 - Introduce **$K-1$** binary discriminant functions, each of which solves a two-class problem of separating points in a particular class **C_k** from points not in that class
 - **One-versus-one classifier**
 - Introduce **$K(K-1)/2$** binary discriminant functions, each one for every possible pair of classes
 - Data are classified according to a **majority vote** amongst the discriminant functions
- The two ways could lead to **some issues**

Example: Three Classes

One-versus-the-rest classifier



One-versus-one classifier



A Single K-Class Discriminant

- Comprise K linear functions, each for a class

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0}$$

- A data point \mathbf{x} is assigned to class \mathcal{C}_k if


$$y_k(\mathbf{x}) > y_j(\mathbf{x}), \forall j \neq k$$

- The decision boundary between class \mathcal{C}_k and class \mathcal{C}_j is given by

$$y_k(\mathbf{x}) = y_j(\mathbf{x})$$

- Which corresponds to a (D-1)-dimensional hyperplane

$$(\mathbf{w}_k - \mathbf{w}_j)^T \mathbf{x} + (w_{k0} - w_{j0}) = 0$$



Least Square for Classification

- Each class C_k is described by its own linear model so that

$$\tilde{\mathbf{w}}_k = [\mathbf{w}_k; w_{k0}] \quad y_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k0}$$

$$\tilde{\mathbf{x}} = [\mathbf{x}; 1] \quad \mathbf{y}(\tilde{\mathbf{x}}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}}$$

- Consider a training dataset with N data points $\{\mathbf{x}_n, \mathbf{t}_n\}$


- Label \mathbf{t} : one-hot encoding (1-of-K binary coding)
 - #Class = 10 (e.g., in digit recognition).
 - One-hot encoding of $t_n = 3$ is $\mathbf{t}_n = [0, 0, 0, 1, 0, 0, 0, 0, 0, 0] \in \{0, 1\}^{10}$

- Least square loss


$$\min_{\tilde{\mathbf{W}}} \sum_{n=1}^N (\mathbf{t}_n - \tilde{\mathbf{W}}^T \tilde{\mathbf{x}}_n)^2$$

Normal equation/(Stochastic) Gradient descent

Least Square for Classification




Magenta line is the decision boundary from least squares




Sensitive to outliers, which lead to large changes in the location of the decision boundary

Least Square for Classification



Linear decision boundaries
could separate classes well



Least square has poor performance

Fisher's Linear Discriminant Analysis

**Linear classification from the viewpoint of
dimensionality reduction**

Essentially, not a discriminant

Main Idea (Binary Classification)

- Project high-dimensional data into a low-dimensional space such that
 - Projected data points from different classes in **low-dim space are separated**
- Project a data point \mathbf{x} to **1 dimension** with a projection vector \mathbf{w} is


$$y = \mathbf{w}^T \mathbf{x}$$

- Goal: maximize the separation of the **projected means** between classes

$$\begin{array}{ll} \mathcal{C}_1 & \mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in \mathcal{C}_1} \mathbf{x}_n \quad m_1 = \mathbf{w}^T \mathbf{m}_1 \\ \mathcal{C}_2 & \mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in \mathcal{C}_2} \mathbf{x}_n \quad m_2 = \mathbf{w}^T \mathbf{m}_2 \end{array}$$

$$\max_{\mathbf{w}} (m_2 - m_1)^2 = (\mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1))^2$$

Issue: Between-Class Overlap



Considerable overlap between classes
when projected onto the 1D line

Fisher's Linear Discriminant

- Further **minimize within-class variance**, thus minimize between-class overlap

$$s_1^2 = \sum_{n \in \mathcal{C}_1} (y_n - m_1)^2 \quad s_2^2 = \sum_{n \in \mathcal{C}_2} (y_n - m_2)^2$$

$$\min_{\mathbf{w}} s_1^2 + s_2^2$$

- Fisher's ratio

$$\frac{\max_{\mathbf{w}} (m_2 - m_1)^2}{\min_{\mathbf{w}} (s_1^2 + w_2^2)} \quad \Rightarrow \quad \max_{\mathbf{w}} \frac{(m_2 - m_1)^2}{(s_1^2 + w_2^2)}$$

Fisher's Linear Discriminant


$$m_k = \mathbf{w}^T \mathbf{m}_k \quad s_k^2 = \sum_{n \in \mathcal{C}_k} (y_n - m_k)^2 \quad y_n = \mathbf{w}^T \mathbf{x}_n$$

Between-class covariance matrix $\mathbf{S}_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T \implies \max_{\mathbf{w}} (m_2 - m_1)^2 = \max_{\mathbf{w}} \mathbf{w}^T \mathbf{S}_B \mathbf{w}$

Total within-class covariance matrix $\mathbf{S}_W = \sum_{k \in \{1,2\}} \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T \implies \min_{\mathbf{w}} (s_1^2 + s_2^2) = \min_{\mathbf{w}} \mathbf{w}^T \mathbf{S}_W \mathbf{w}$

$$\max_{\mathbf{w}} \frac{(m_2 - m_1)^2}{(s_1^2 + s_2^2)} \implies \max_{\mathbf{w}} \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

Differentiate w.r.t \mathbf{w} and set it to be $\mathbf{0}$



$$(\mathbf{w}^T \mathbf{S}_B \mathbf{w}) \mathbf{S}_W \mathbf{w} = (\mathbf{w}^T \mathbf{S}_W \mathbf{w}) \mathbf{S}_B \mathbf{w}$$

$$\mathbf{S}_B \mathbf{w} = (\mathbf{m}_2 - \mathbf{m}_1)((\mathbf{m}_2 - \mathbf{m}_1)^T \mathbf{w}) \propto \mathbf{m}_2 - \mathbf{m}_1 \implies \mathbf{w} \propto \mathbf{S}_W^{-1} (\mathbf{m}_2 - \mathbf{m}_1)$$

Fisher's Linear Discriminant

$$\mathbf{w} \propto \mathbf{S}_{\mathbf{w}}^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$


Classification $y(\mathbf{x}_t) = \mathbf{w}^T \mathbf{x}_t$

$$= (\mathbf{m}_2 - \mathbf{m}_1)^T \mathbf{S}_{\mathbf{w}}^{-1} \mathbf{x}_t \quad \mathcal{C}_1$$

$$\geq y_0 \quad \text{Need to determine } y_0$$

It is essentially not a discriminant

A Two-Class Example



Between-class overlap is significantly reduced
Within-class data points are close (small variance)

Generalize to Multi-Classes ($K < D$)

$$\mathbf{w} \in \mathbb{R}^D \quad y_n = \mathbf{w}^T \mathbf{x}_n \quad \Longrightarrow \quad \mathbf{y}_n = \mathbf{W}^T \mathbf{x}_n \quad \mathbf{W} \in \mathbb{R}^{D \times d}$$

Per-class mean $\mathbf{m}_k = \frac{1}{N_k} \sum_{n \in \mathcal{C}_k} \mathbf{x}_n$ $\mathbf{m} = \frac{1}{N} \sum_{n=1}^N \mathbf{x}_n = \frac{1}{N} \sum_{k=1}^K N_k \mathbf{m}_k$ All-class mean


Between-class covariance matrix $\mathbf{S}_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T \quad \Longrightarrow \quad \mathbf{S}_B = \sum_{k=1}^K N_k (\mathbf{m}_k - \mathbf{m})(\mathbf{m}_k - \mathbf{m})^T$

Total within-class covariance matrix $\mathbf{S}_W = \sum_{k \in \{1,2\}} \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T \quad \Longrightarrow \quad \mathbf{S}_W = \sum_{k=1}^K \sum_{n \in \mathcal{C}_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T$


Fisher's ratio $\max_{\mathbf{w}} \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}} \quad \Longrightarrow \quad \max_{\mathbf{w}} \text{Tr}\{(\mathbf{W}^T \mathbf{S}_W \mathbf{W})^{-1} (\mathbf{W}^T \mathbf{S}_B \mathbf{W})\}$

To create a discriminant, we model a Gaussian distribution over the D -dim data \mathbf{x} for each class k

10-Classes MNIST Digit Classification



d=2



d=3

Perceptron Algorithm

(Only for Binary Classification)


Perceptron Algorithm

- Another example of a linear discriminant function
 - An important place in the history of pattern recognition (**Rosenblatt**, 1962)
- Data point **x** is first transformed using a nonlinear transformation **ϕ** to give a feature vector **$\phi(\mathbf{x})$** ;
- Then apply a nonlinear activation function (step function) **f** to *classify data*

$$y = f(\mathbf{w}^T \phi(\mathbf{x})) \quad f(a) = \begin{cases} +1, & \text{if } a > 0; \\ -1, & \text{if } a < 0. \end{cases}$$

$$\mathbf{w}^T \phi(\mathbf{x}) > 0 \quad +1$$

$$\mathbf{w}^T \phi(\mathbf{x}) < 0 \quad -1$$



Perceptron Algorithm

- Binary classification: label $t \in \{-1, +1\}$

$$\begin{aligned} \mathbf{x}_n \in \mathcal{C}_1 : t_n = +1 &\implies \mathbf{w}^T \phi(\mathbf{x}_n) > 0 \\ \mathbf{x}_n \in \mathcal{C}_2 : t_n = -1 &\implies \mathbf{w}^T \phi(\mathbf{x}_n) < 0 \end{aligned} \implies \mathbf{w}^T \phi(\mathbf{x}_n) \cdot t_n > 0$$

- The perceptron has **zero error** with any data point **correctly classified**

$$t_n = +1 \text{ (OR } -1) \quad y_n = f(\mathbf{w}^T \phi(\mathbf{x}_n)) = +1 \text{ (OR } -1)$$

- Whereas a misclassified data \mathbf{x}_n it incurs an error

$$-\mathbf{w}^T \phi(\mathbf{x}_n) t_n$$

- The total error of perceptron

$$E_P(\mathbf{w}) = \sum_{n \in \mathcal{M}} -\mathbf{w}^T \phi(\mathbf{x}_n) t_n \quad \mathbf{M} \text{ is the set of all misclassified data}$$

Solution: Stochastic Gradient Descent

- Per sample gradient descent
 - For a misclassified data point \mathbf{x}_n

$$E(\mathbf{w}; \mathbf{x}_n) = -\mathbf{w}^T \phi(\mathbf{x}_n) t_n \implies \nabla_{\mathbf{w}} E(\mathbf{w}; \mathbf{x}_n) = -\phi(\mathbf{x}_n) t_n$$
$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \nabla_{\mathbf{w}^{(t)}} E(\mathbf{w}^{(t)}; \mathbf{x}_n) = \mathbf{w}^{(t)} + \phi(\mathbf{x}_n) t_n$$

- Simple interpretation
 - If a data is correctly classified, then the weight vector remains unchanged
 - If it is incorrectly classified, there is a penalty $|\phi(\mathbf{x}_n)|$
- Error from a data point is reduced with a single update

$$-(\mathbf{w}^{(t+1)})^T \phi(\mathbf{x}_n) t_n = -(\mathbf{w}^{(t)})^T \phi(\mathbf{x}_n) t_n - \|\phi(\mathbf{x}_n) t_n\|^2 < -(\mathbf{w}^{(t)})^T \phi(\mathbf{x}_n) t_n$$

Perceptron Convergence & Correctness

- Convergence
 - If training data is **linearly separable**, then perceptron learning is guaranteed to find **an exact solution** in a **finite number of iterations**

- Correctness

- Assume the length of all data points is bounded by D , i.e.,

$$\|\mathbf{x}_n\|_2 \leq D, \forall n$$


- Assume there exists unit length \mathbf{w} and some $\gamma > 0$ such that

$$\mathbf{w}^T \phi(\mathbf{x}_n) \cdot t_n > \gamma, \forall n$$


- The **total number of mistakes** the perceptron algorithm makes is **at most**

$$(D/\gamma)^2$$

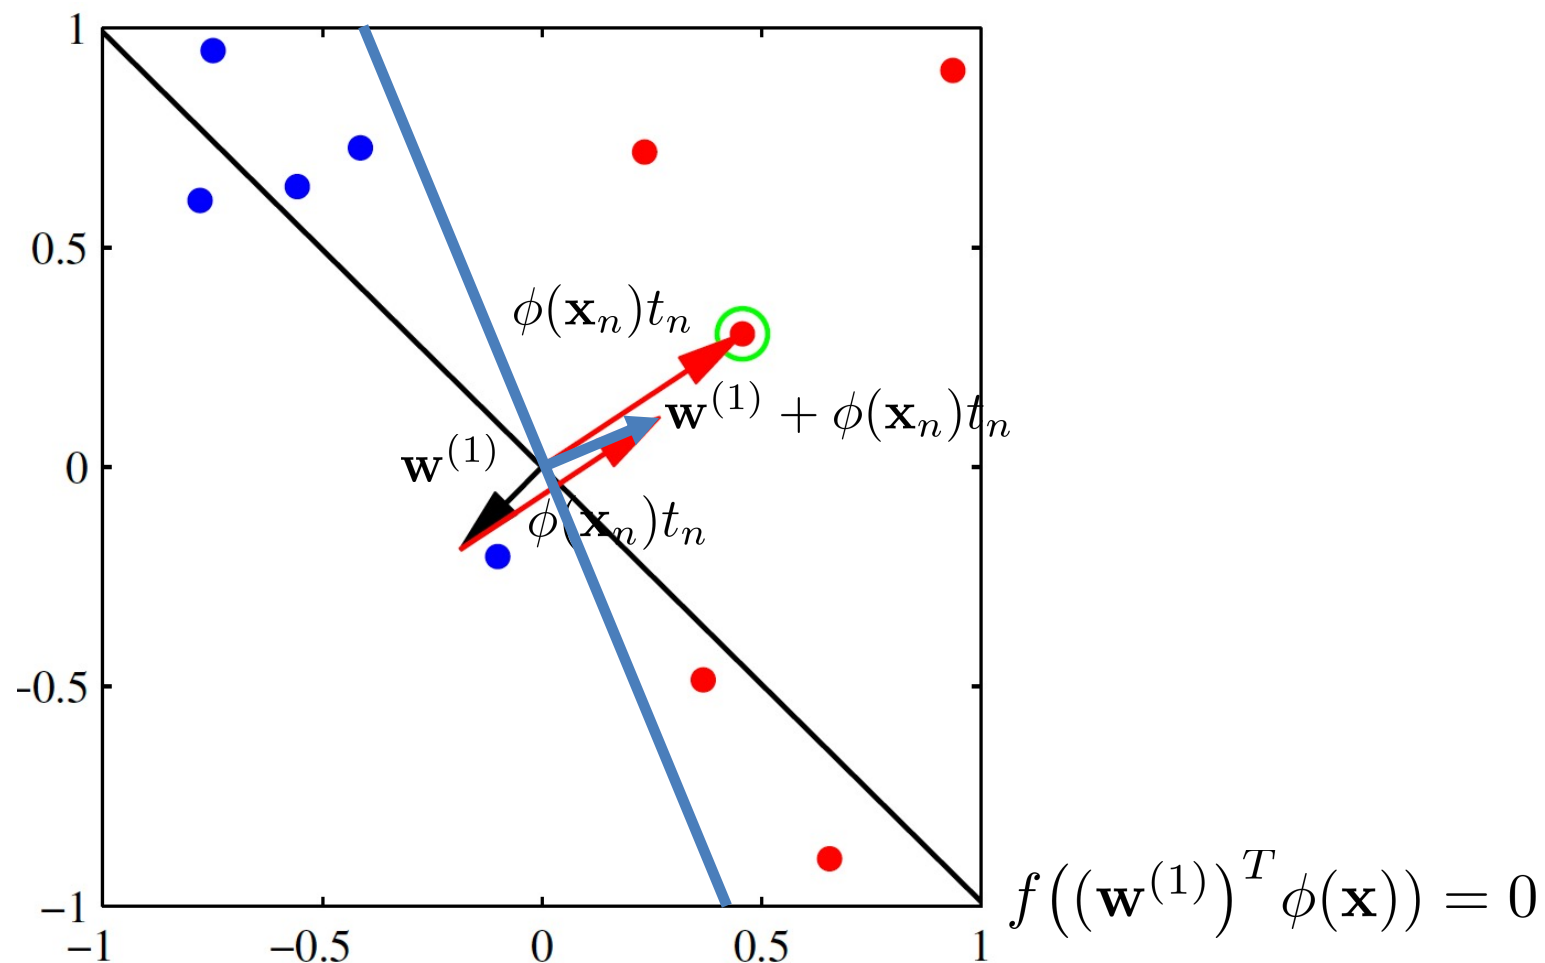
Perceptron Algorithm: Example




Perceptron Algorithm: Example



Perceptron Algorithm: Example



Perceptron Algorithm: Example



Perceptron: Pros & Cons

- Pros
 - Easy to implement
 - Time/memory efficient
 - Guaranteed performance when data points are linearly separable
- Cons
 - **Sensitive** to initialized parameter vector
 - Only applicable to **binary classification**
 - **NEVER converge** when data points are **not linearly separable**